

Homework 3 Suglia

Elena Suglia

11/17/2019

Question 1

Using the “litterbags.csv” data set, create and run a model in JAGS that corresponds to this one:

```
lmer(N_min_rate~Celastrus + (1|Plot), data=litterbags)
```

This data set contains nitrogen mineralization rates at 7 different plots where the treatment was whether or not the invasive liana *Celastrus orbiculatus* was present or not.

In your answer, please include the model code and the means and standard deviations of the intercept, slope, group-level (i.e. among-plot) variance and individual-level (within-plot) variance. Briefly report what you did to check model convergence.

Model code in text file

```
model {  
  
  # Likelihood  
  for (i in 1:n) {  
    y[i] ~ dnorm(y.hat[i], tau.y)  
    y.hat[i] <- a[Plot[i]] + b*x[i] # regression with intercept for each plot  
  }  
  
  # group-level model  
  for (j in 1:7) { # 7 plots  
    a[j] ~ dnorm(mu.a, tau.a) # separate intercepts for every plot, no pooling of information  
  }  
  
  #priors  
  mu.a ~ dnorm(0, 0.1) # overall mean  
  sigma.a ~ dnorm(0, 0.5)T(0,) # weakly informative prior on the group-level standard deviation  
  tau.a <- pow(sigma.a, -2) # group-level precision  
  
  # individual-level priors  
  b ~ dnorm(0, 0.1) # prior for slope  
  sigma.y ~ dnorm(0, 0.1)T(0,) # weakly informative prior on the individual-level sd  
  tau.y <- pow(sigma.y, -2) # individual-level precision  
}
```

Read in data

```
d = read_csv("litterbags.csv")
```

Define data needed for model

```
length(unique(d$Plot)) # 7 plots
d2.data <- list(n=nrow(d), y=d$N_min_rate, x=d$Celastrus,
plot.index=d$Plot, n.plots=length(unique(d$Plot)))
```

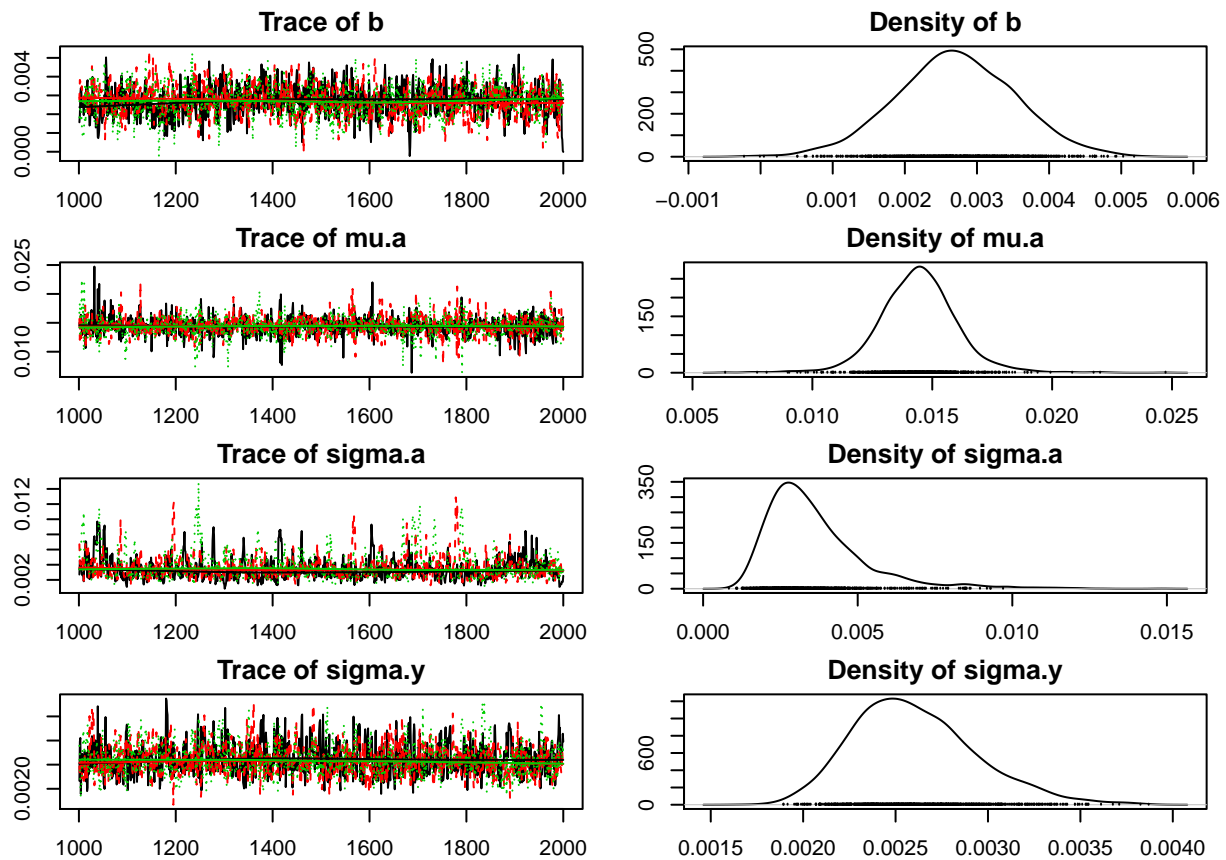
Have jags generate starting values

```
d2.inits <- list(list(sigma.a = 1, sigma.y=2), list(sigma.a = 2,
sigma.y=1), list(sigma.a = 5, sigma.y=0.4))
d2 <- jags.model("litterbags_jags_hw3.txt", data=d2.data, inits = d2.inits,
n.chains=3, n.adapt=1000)
```


Do a trial run with some parameters, and check convergence

```
d2.samp <- coda.samples(d2, c("b", "sigma.y", "mu.a", "sigma.a"), n.iter=1000)
summary(d2.samp)
```

```
##
## Iterations = 1001:2000
## Thinning interval = 1
## Number of chains = 3
## Sample size per chain = 1000
##
## 1. Empirical mean and standard deviation for each variable,
##    plus standard error of the mean:
##
##              Mean          SD Naive SE Time-series SE
## b              0.002702 0.0008266 1.509e-05      2.628e-05
## mu.a           0.014418 0.0016406 2.995e-05      3.550e-05
## sigma.a        0.003610 0.0016476 3.008e-05      7.688e-05
## sigma.y        0.002603 0.0003317 6.055e-06      9.707e-06
##
## 2. Quantiles for each variable:
##
##              2.5%      25%      50%      75%      97.5%
## b              0.001023 0.002165 0.002697 0.003268 0.004338
## mu.a           0.011286 0.013450 0.014412 0.015348 0.017777
## sigma.a        0.001607 0.002520 0.003213 0.004243 0.008381
## sigma.y        0.002041 0.002364 0.002569 0.002801 0.003336
par(mar=rep(2, 4))
plot(d2.samp)
```



```
acfplot(d2.samp)
```



```
gelman.diag(d2.samp)
```

```
## Potential scale reduction factors:
##
##           Point est. Upper C.I.
## b           1.00         1.00
## mu.a         1.00         1.00
## sigma.a       1.03         1.07
## sigma.y       1.00         1.01
##
## Multivariate psrf
##
## 1.01
```

Add more iterations and include all parameters of interest

```
d2.samp <- coda.samples(d2, c("a", "b", "sigma.y", "mu.a", "sigma.a", "y.hat"), n.iter=5000, thin=5)
```

Look at the results

```
d2.summ <- summary(d2.samp)
d2.stats <- as.data.frame(d2.summ$statistics)
head(d2.stats, 100)
```

##		Mean	SD	Naive SE	Time-series SE
##	a[1]	0.018094452	0.0011581855	2.114548e-05	2.114688e-05
##	a[2]	0.012893579	0.0010864376	1.983555e-05	1.981748e-05
##	a[3]	0.017677102	0.0012783733	2.333980e-05	2.273521e-05
##	a[4]	0.012115220	0.0011581635	2.114508e-05	2.081226e-05
##	a[5]	0.014956385	0.0011040788	2.015763e-05	2.144425e-05
##	a[6]	0.012622608	0.0010816347	1.974786e-05	2.019552e-05
##	a[7]	0.012773676	0.0011091395	2.025002e-05	1.948926e-05
##	b	0.002665172	0.0008477275	1.547732e-05	1.615134e-05
##	mu.a	0.014433530	0.0015623524	2.852452e-05	2.852923e-05
##	sigma.a	0.003491457	0.0015403687	2.812316e-05	3.243161e-05
##	sigma.y	0.002616677	0.0003394741	6.197921e-06	6.332734e-06
##	y.hat[1]	0.020759624	0.0011665434	2.129807e-05	2.063974e-05
##	y.hat[2]	0.020759624	0.0011665434	2.129807e-05	2.063974e-05
##	y.hat[3]	0.020759624	0.0011665434	2.129807e-05	2.063974e-05
##	y.hat[4]	0.018094452	0.0011581855	2.114548e-05	2.114688e-05
##	y.hat[5]	0.018094452	0.0011581855	2.114548e-05	2.114688e-05
##	y.hat[6]	0.018094452	0.0011581855	2.114548e-05	2.114688e-05
##	y.hat[7]	0.015558751	0.0010988376	2.006194e-05	1.976811e-05
##	y.hat[8]	0.015558751	0.0010988376	2.006194e-05	1.976811e-05
##	y.hat[9]	0.015558751	0.0010988376	2.006194e-05	1.976811e-05
##	y.hat[10]	0.012893579	0.0010864376	1.983555e-05	1.981748e-05
##	y.hat[11]	0.012893579	0.0010864376	1.983555e-05	1.981748e-05
##	y.hat[12]	0.012893579	0.0010864376	1.983555e-05	1.981748e-05
##	y.hat[13]	0.020342275	0.0012156266	2.219420e-05	2.218613e-05
##	y.hat[14]	0.020342275	0.0012156266	2.219420e-05	2.218613e-05
##	y.hat[15]	0.020342275	0.0012156266	2.219420e-05	2.218613e-05

```
## y.hat[16] 0.017677102 0.0012783733 2.333980e-05 2.273521e-05
## y.hat[17] 0.017677102 0.0012783733 2.333980e-05 2.273521e-05
## y.hat[18] 0.014780392 0.0011507925 2.101050e-05 2.164535e-05
## y.hat[19] 0.014780392 0.0011507925 2.101050e-05 2.164535e-05
## y.hat[20] 0.014780392 0.0011507925 2.101050e-05 2.164535e-05
## y.hat[21] 0.012115220 0.0011581635 2.114508e-05 2.081226e-05
## y.hat[22] 0.012115220 0.0011581635 2.114508e-05 2.081226e-05
## y.hat[23] 0.012115220 0.0011581635 2.114508e-05 2.081226e-05
## y.hat[24] 0.017621557 0.0010966608 2.002220e-05 2.002001e-05
## y.hat[25] 0.017621557 0.0010966608 2.002220e-05 2.002001e-05
## y.hat[26] 0.017621557 0.0010966608 2.002220e-05 2.002001e-05
## y.hat[27] 0.014956385 0.0011040788 2.015763e-05 2.144425e-05
## y.hat[28] 0.014956385 0.0011040788 2.015763e-05 2.144425e-05
## y.hat[29] 0.014956385 0.0011040788 2.015763e-05 2.144425e-05
## y.hat[30] 0.015287781 0.0011196661 2.044221e-05 1.997699e-05
## y.hat[31] 0.015287781 0.0011196661 2.044221e-05 1.997699e-05
## y.hat[32] 0.015287781 0.0011196661 2.044221e-05 1.997699e-05
## y.hat[33] 0.012622608 0.0010816347 1.974786e-05 2.019552e-05
## y.hat[34] 0.012622608 0.0010816347 1.974786e-05 2.019552e-05
## y.hat[35] 0.012622608 0.0010816347 1.974786e-05 2.019552e-05
## y.hat[36] 0.015438848 0.0011179322 2.041056e-05 2.039300e-05
## y.hat[37] 0.015438848 0.0011179322 2.041056e-05 2.039300e-05
## y.hat[38] 0.015438848 0.0011179322 2.041056e-05 2.039300e-05
## y.hat[39] 0.012773676 0.0011091395 2.025002e-05 1.948926e-05
## y.hat[40] 0.012773676 0.0011091395 2.025002e-05 1.948926e-05
## y.hat[41] 0.012773676 0.0011091395 2.025002e-05 1.948926e-05
```

```
y.hat.rows <- grep("y.hat", rownames(d2.stats))
y.hat <- d2.stats$Mean[y.hat.rows] # To get fitted values of the model, look at the stats and pull out
resids <- d$N_min_rate - y.hat
# for convenience, we can then put those values into our data frame and examine them
d <- cbind(d, y.hat = y.hat, resid = resids)
```

Some summary statistics:

The means and standard deviations of:

- slope = b
- intercept = mu.a
- group-level (i.e. among-plot) variance = sigma.a
- individual-level (within-plot) variance = sigma.y


```
##           Mean          SD
## a[1]      0.018094452 0.0011581855
## a[2]      0.012893579 0.0010864376
## a[3]      0.017677102 0.0012783733
## a[4]      0.012115220 0.0011581635
## a[5]      0.014956385 0.0011040788
## a[6]      0.012622608 0.0010816347
## a[7]      0.012773676 0.0011091395
## b         0.002665172 0.0008477275
## mu.a      0.014433530 0.0015623524
## sigma.a   0.003491457 0.0015403687
## sigma.y   0.002616677 0.0003394741
```

```
## y.hat[1] 0.020759624 0.0011665434
## y.hat[2] 0.020759624 0.0011665434
## y.hat[3] 0.020759624 0.0011665434
## y.hat[4] 0.018094452 0.0011581855
## y.hat[5] 0.018094452 0.0011581855
## y.hat[6] 0.018094452 0.0011581855
## y.hat[7] 0.015558751 0.0010988376
## y.hat[8] 0.015558751 0.0010988376
## y.hat[9] 0.015558751 0.0010988376
## y.hat[10] 0.012893579 0.0010864376
## y.hat[11] 0.012893579 0.0010864376
## y.hat[12] 0.012893579 0.0010864376
## y.hat[13] 0.020342275 0.0012156266
## y.hat[14] 0.020342275 0.0012156266
## y.hat[15] 0.020342275 0.0012156266
## y.hat[16] 0.017677102 0.0012783733
## y.hat[17] 0.017677102 0.0012783733
## y.hat[18] 0.014780392 0.0011507925
## y.hat[19] 0.014780392 0.0011507925
## y.hat[20] 0.014780392 0.0011507925
## y.hat[21] 0.012115220 0.0011581635
## y.hat[22] 0.012115220 0.0011581635
## y.hat[23] 0.012115220 0.0011581635
## y.hat[24] 0.017621557 0.0010966608
## y.hat[25] 0.017621557 0.0010966608
## y.hat[26] 0.017621557 0.0010966608
## y.hat[27] 0.014956385 0.0011040788
## y.hat[28] 0.014956385 0.0011040788
## y.hat[29] 0.014956385 0.0011040788
## y.hat[30] 0.015287781 0.0011196661
## y.hat[31] 0.015287781 0.0011196661
## y.hat[32] 0.015287781 0.0011196661
## y.hat[33] 0.012622608 0.0010816347
## y.hat[34] 0.012622608 0.0010816347
## y.hat[35] 0.012622608 0.0010816347
## y.hat[36] 0.015438848 0.0011179322
## y.hat[37] 0.015438848 0.0011179322
## y.hat[38] 0.015438848 0.0011179322
## y.hat[39] 0.012773676 0.0011091395
## y.hat[40] 0.012773676 0.0011091395
## y.hat[41] 0.012773676 0.0011091395
```


See how well the model ran

```
qqnorm(d$resid)
```

Normal Q-Q Plot



```
boxplot(resid~Plot, data=d)
```



Now that random intercepts for county are included in the model, does it look like the residuals diff

Look at DIC

```
d2.DIC <- dic.samples(d2, n.iter=5000, thin=5)
d2.DIC

## Mean deviance: -373.2
## penalty 8.931
## Penalized deviance: -364.2
```

For comparison, let's quickly fit the complete-pooling model that has no county effects.

```
#d1.data <- list(n=nrow(d), y=dNminrate, x = dCelastrus) # for this simple model we can let jags generate
starting values #d1 <- jags.model("d_completenesspool_jags.txt", data=d1.data, n.chains=3, n.adapt=1000)
#d1.samp <- coda.samples(d1, c("a", "b", "sigma.y"), n.iter=1000) #d1.DIC <- dic.samples(d1, n.iter=1000)
```

Question 2

Fit a hierarchical (multilevel) model in JAGS using the data sets “immunity.csv” and “patient_age.csv”. The model should include “immune.level” as the response variable, “time” as an individual-level predictor, and “age” as a group-level predictor.

These data are repeated measures on individual patients, so each patient is a “group” in the data set. Each row in the data set “immunity.csv” is one observation on one patient. Each row in the data set “patient_age.csv” is the age of each patient at time the study began – so this is a group level predictor, with one row of data per patient. In other words, here “age” is analogous to county-level bedrock uranium content in the Gelman & Hill radon example.

Hints: You can use the column “patient” in “immunity.csv” to index the random intercept for patient. You can then use the column “age” in “patient_age.csv” in the group-level regression that explains some of the variation in the random intercepts. There is a model like this on page 361 of Gelman & Hill.

In your answer, include the model you created. Also report the means and standard deviations of: - the slopes of the individual-data-point-level and “group”-level (i.e. patient-level) regressions - the individual-data-point-level and “group”-level variance parameters